Probability and Statistics for Computer Science

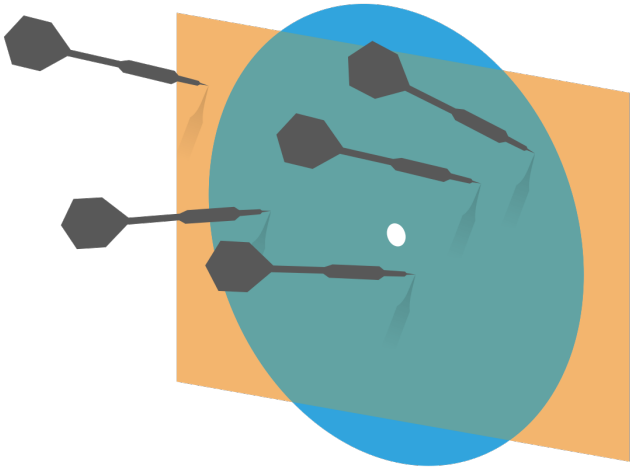

Credit: wikipedia

"Correlation is not Causation"
but Correlation is so beautiful!

Hongye Liu, Teaching Assistant Prof, CS361, UIUC, 8.29.2023

Last time

- Mean
- Standard deviation
- Variance
- Standardizing data

$\{x_i\}$ $i \in [1, 8]$

$\{x_i\} = 1, 2, 3, 4, 5, 6, 7, 12$

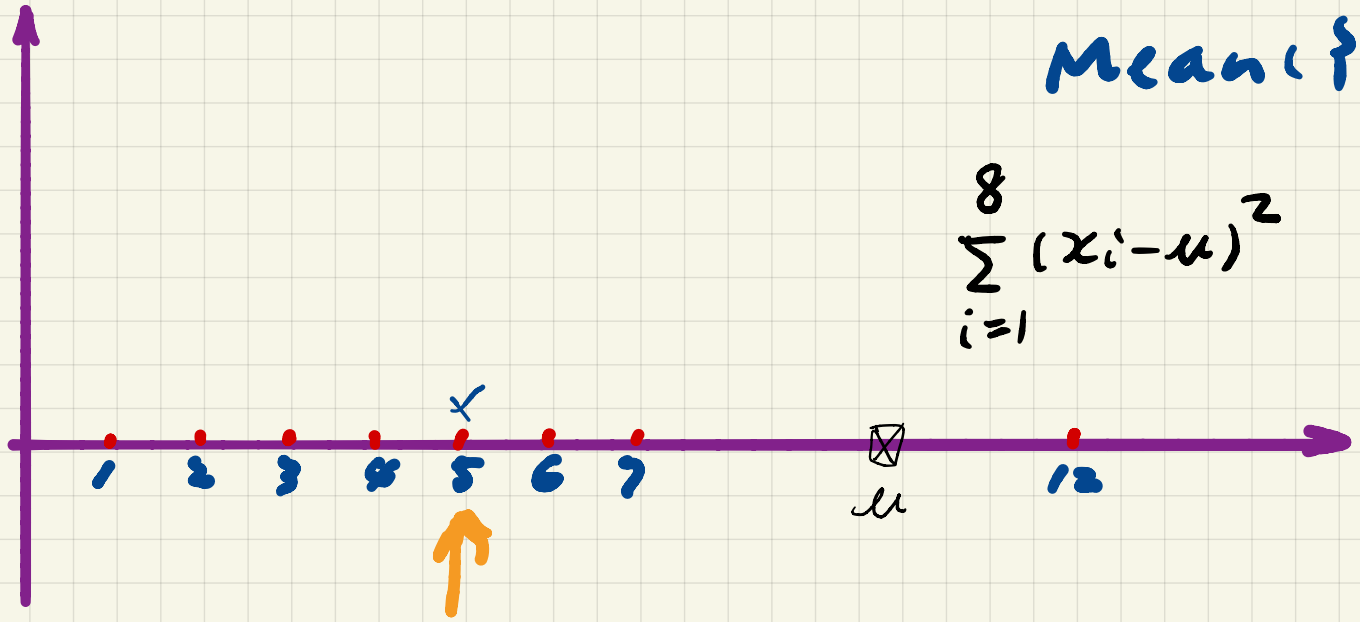

$\text{Mean}(\{x_i\}) = 5$

$$\sum_{i=1}^{8} (x_i - \mu)^2$$

$$\underset{\mu}{\operatorname{argmin}} \left(\sum_i (x_i - \mu)^2 \right)$$

$\{x_i\}$ $\quad i \in [1, N]$

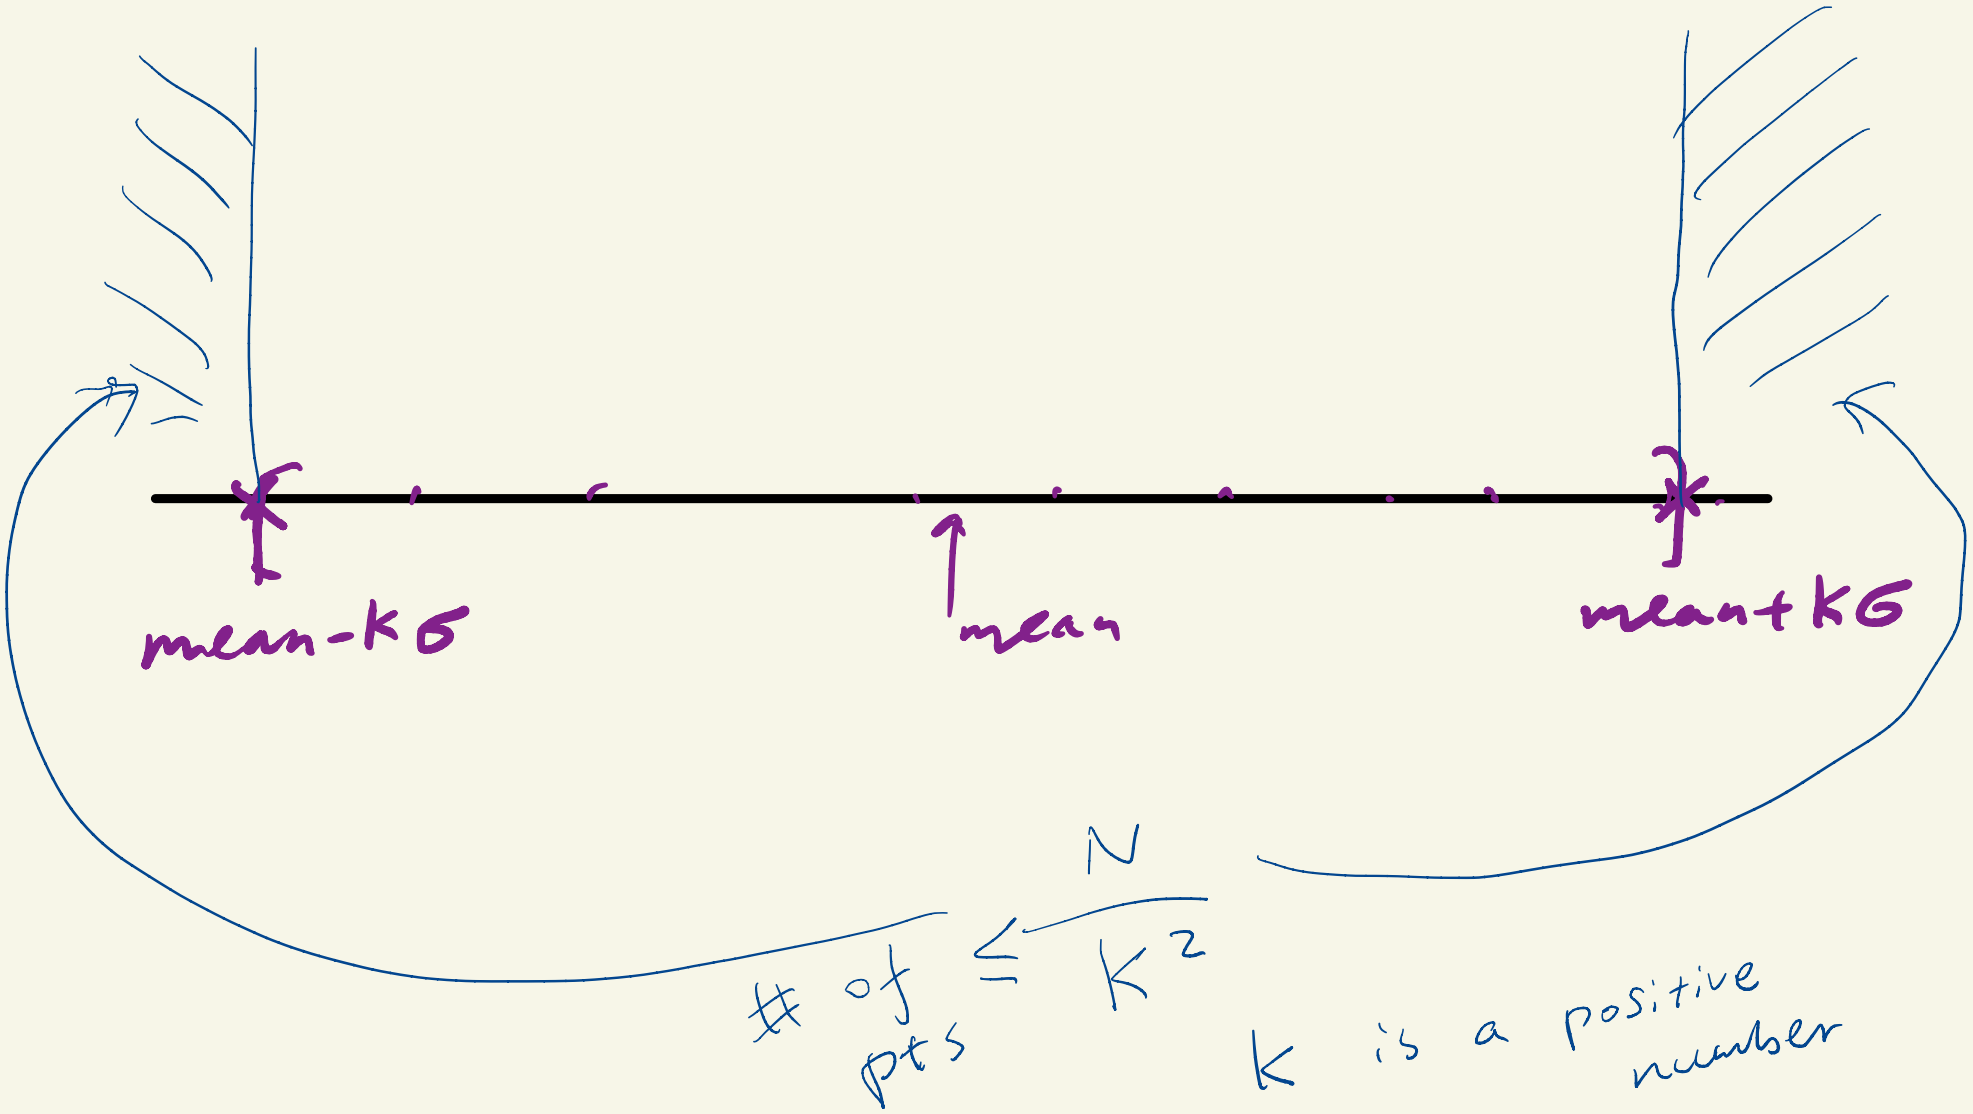

$$\#\text{ of pts} \leq \frac{N}{k^2}$$

k is a positive number

Q: Estimate the range of data in standard coordinates

- The interval [-5, 5] covers x% data, choose the closest estimate of x.

A. 80%

B. 99

C. 96

D. 96%

$$\widehat{x_i} = \frac{x_i - mean(\{x_i\})}{std(\{x_i\})}$$

Objectives

- Median, Percentile, Mode, IQR,
- Scatter plots for relationships
- Correlation Coefficient
- Other Visualization for *relationships*
  Heatmap, 3D bar, Time series plots,

Median

- We first sort the data set $\{x_i\}$
- Then *if* the number of items N is odd

  median = middle item's value

  if the number of items N is even

  median = mean of middle 2 items' values

Properties of Median

- Scaling data scales the median

$$median(\{k \cdot x_i\}) = k \cdot median(\{x_i\})$$

- Translating data translates the median

$$median(\{x_i + c\}) = median(\{x_i\}) + c$$

Percentile

- k^{th} percentile is the value relative to which k% of the data items have smaller or equal numbers
- Median is roughly the 50^{th} percentile

Interquartile range

- iqr = (75th percentile) - (25th percentile)
- Scaling data scales the interquartile range

$$iqr(\{k \cdot x_i\}) = |k| \cdot iqr(\{x_i\})$$

- Translating data does **NOT** change the interquartile range

$$iqr(\{x_i + c\}) = iqr(\{x_i\})$$

Box plots

- Boxplots
  - Simpler than histogram
  - Good for outliers
  - Easier to use for comparison

Vehicle death by region

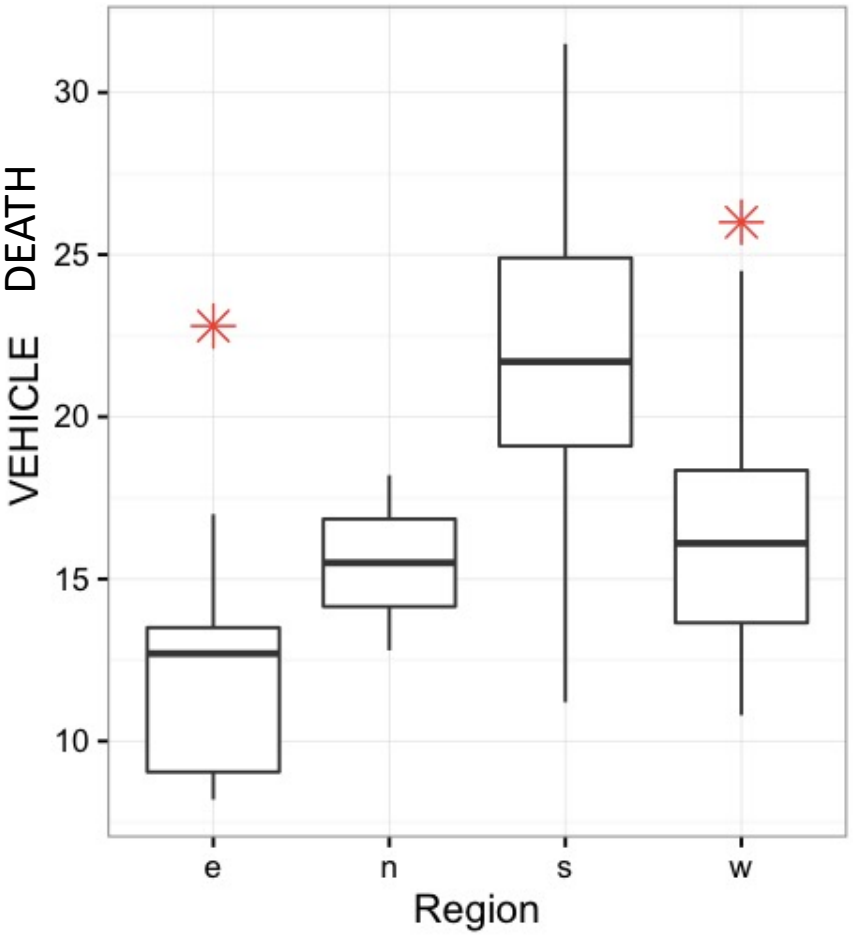

Data from
https://www2.stetson.edu/~jrasp/data.htm

Boxplots details, outliers

- How to define outliers? (the default)

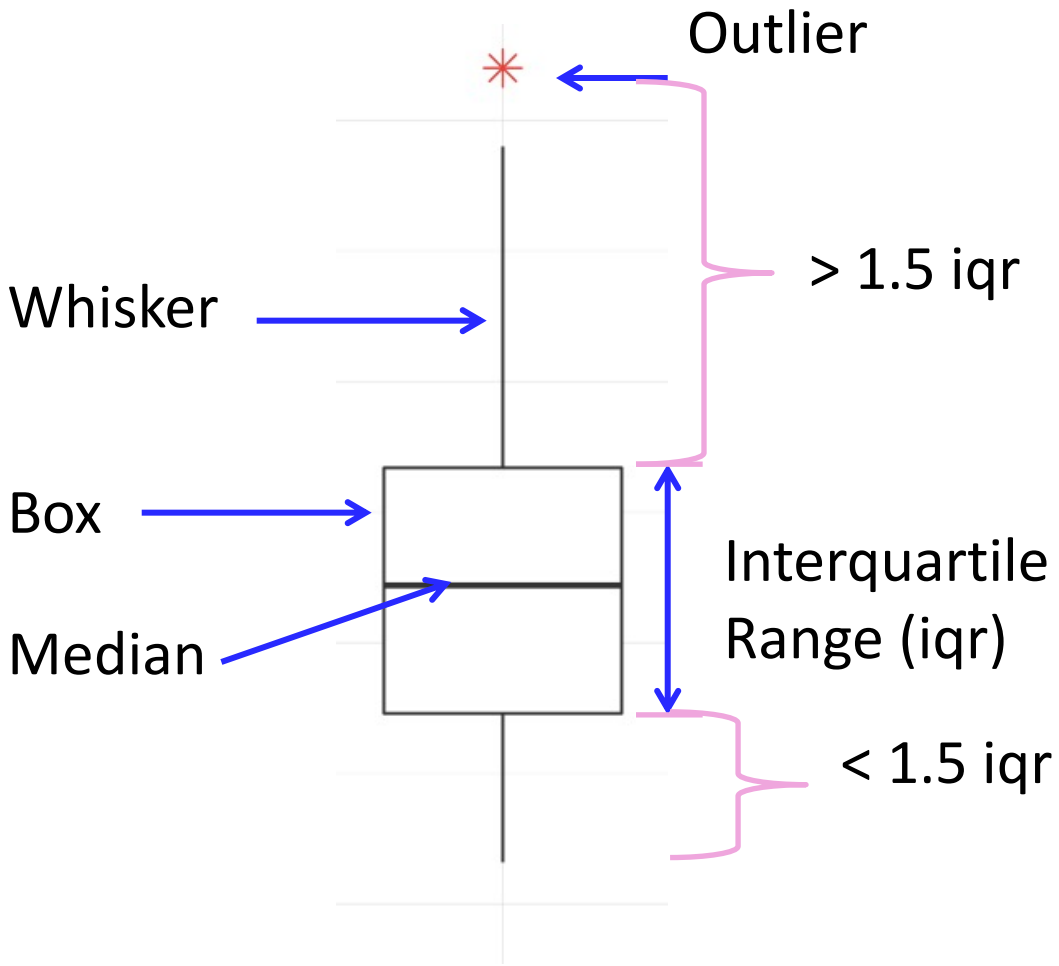

Sensitivity of summary statistics to outliers

- mean and standard deviation are very sensitive to outliers
- median and interquartile range are not sensitive to outliers

Modes

- Modes are peaks in a histogram
- If there are more than 1 mode, we should be curious as to why

Multiple modes

- We have seen the "iris" data which looks to have several peaks

Data: "iris" in R

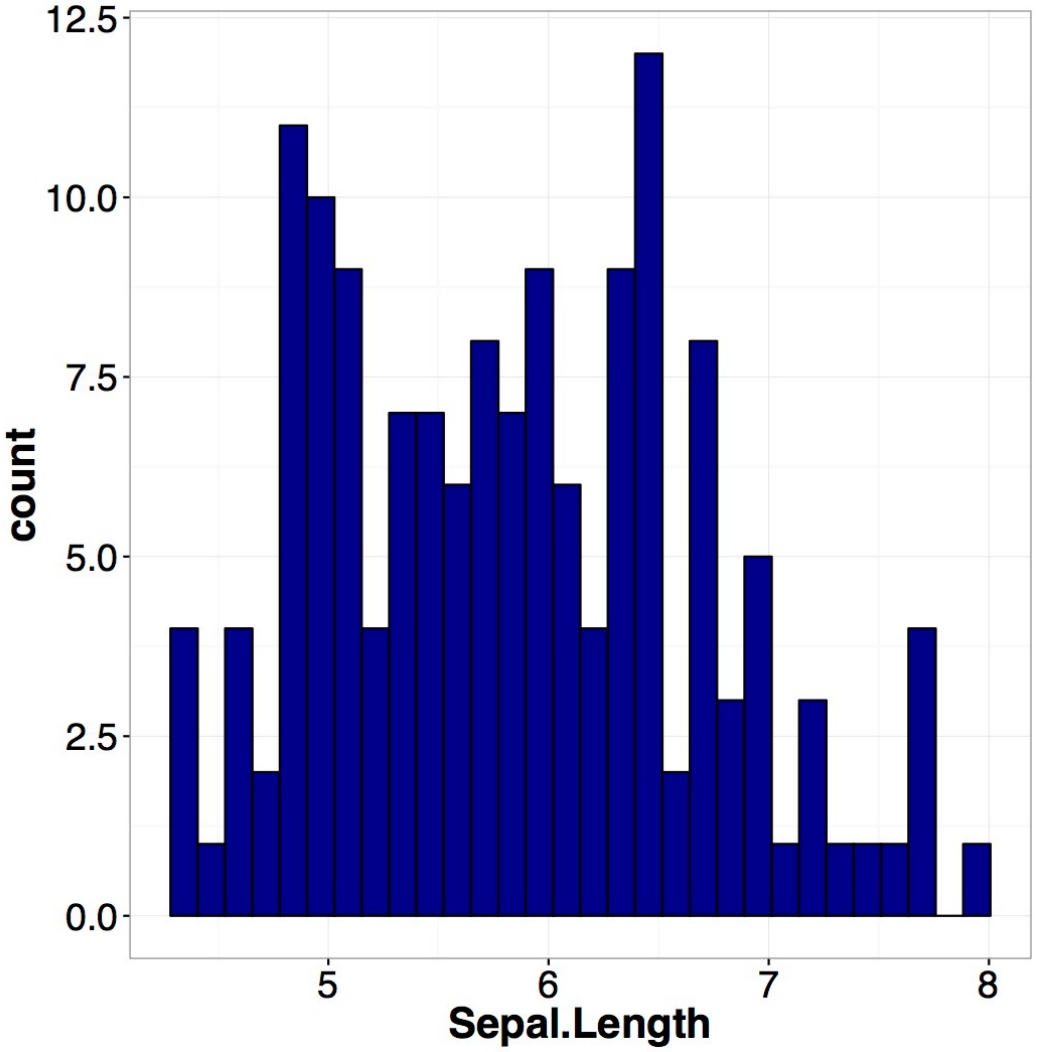

Example Bi-modes distribution

- Modes may indicate multiple populations

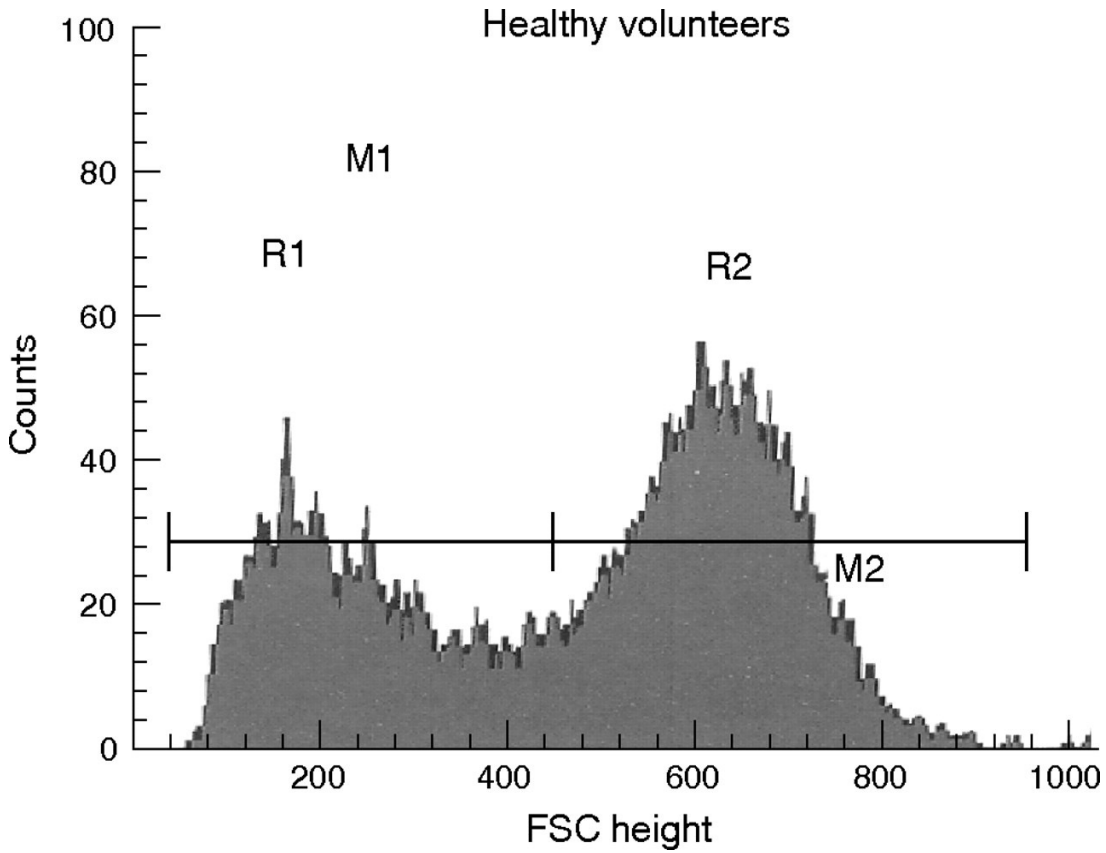

Data: Erythrocyte cells in healthy humans

Piagnerelli, JCP 2007

Tails and Skews

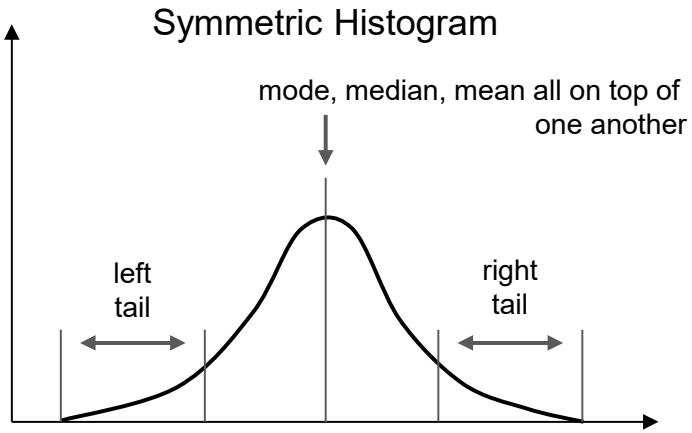

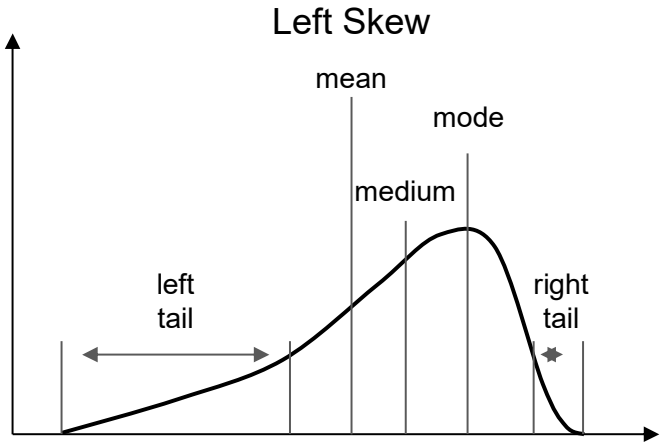

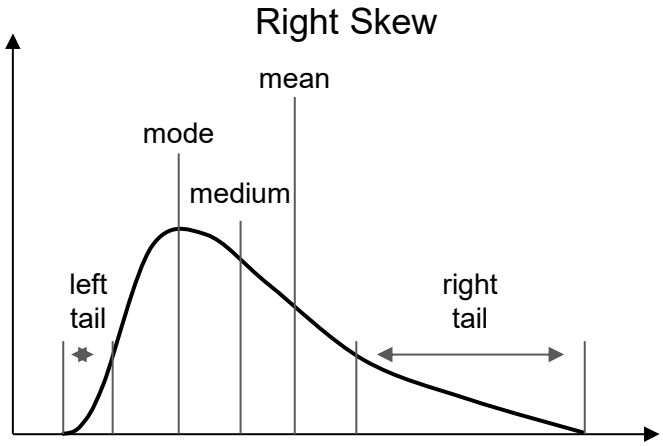

Credit: Prof.Forsyth

Relationship between data features

- Example: Does the weight of people relate to their height?

IDNO	BODYFAT	DENSITY	AGE	WEIGHT	HEIGHT
1	12.6	1.0708	23	154.25	67.75
2	6.9	1.0853	22	173.25	72.25
3	24.6	1.0414	22	154.00	66.25
4	10.9	1.0751	26	184.75	72.25
5	27.8	1.0340	24	184.25	71.25
6	20.6	1.0502	24	210.25	74.75
7	19.0	1.0549	26	181.00	69.75
8	12.8	1.0704	25	176.00	72.50
9	5.1	1.0900	25	191.00	74.00
10	12.0	1.0722	23	198.25	73.50

- x : HIGHT, y: WEIGHT

Scatter plot

- Body Fat data set

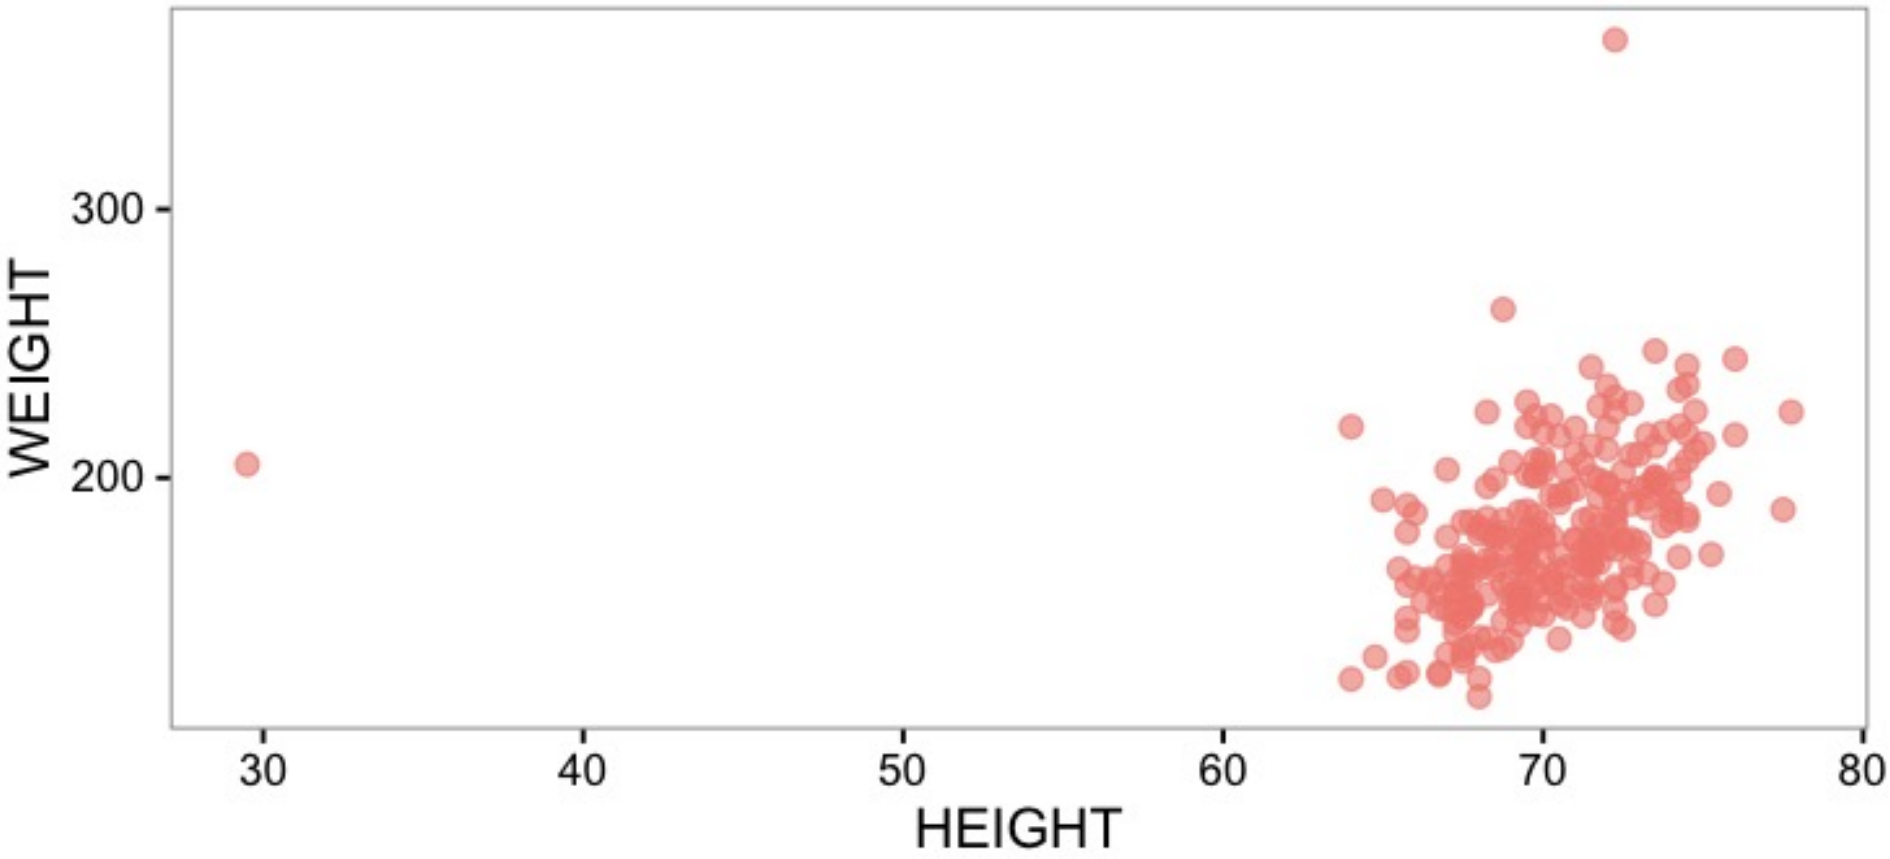

Scatter plot

- Scatter plot with density

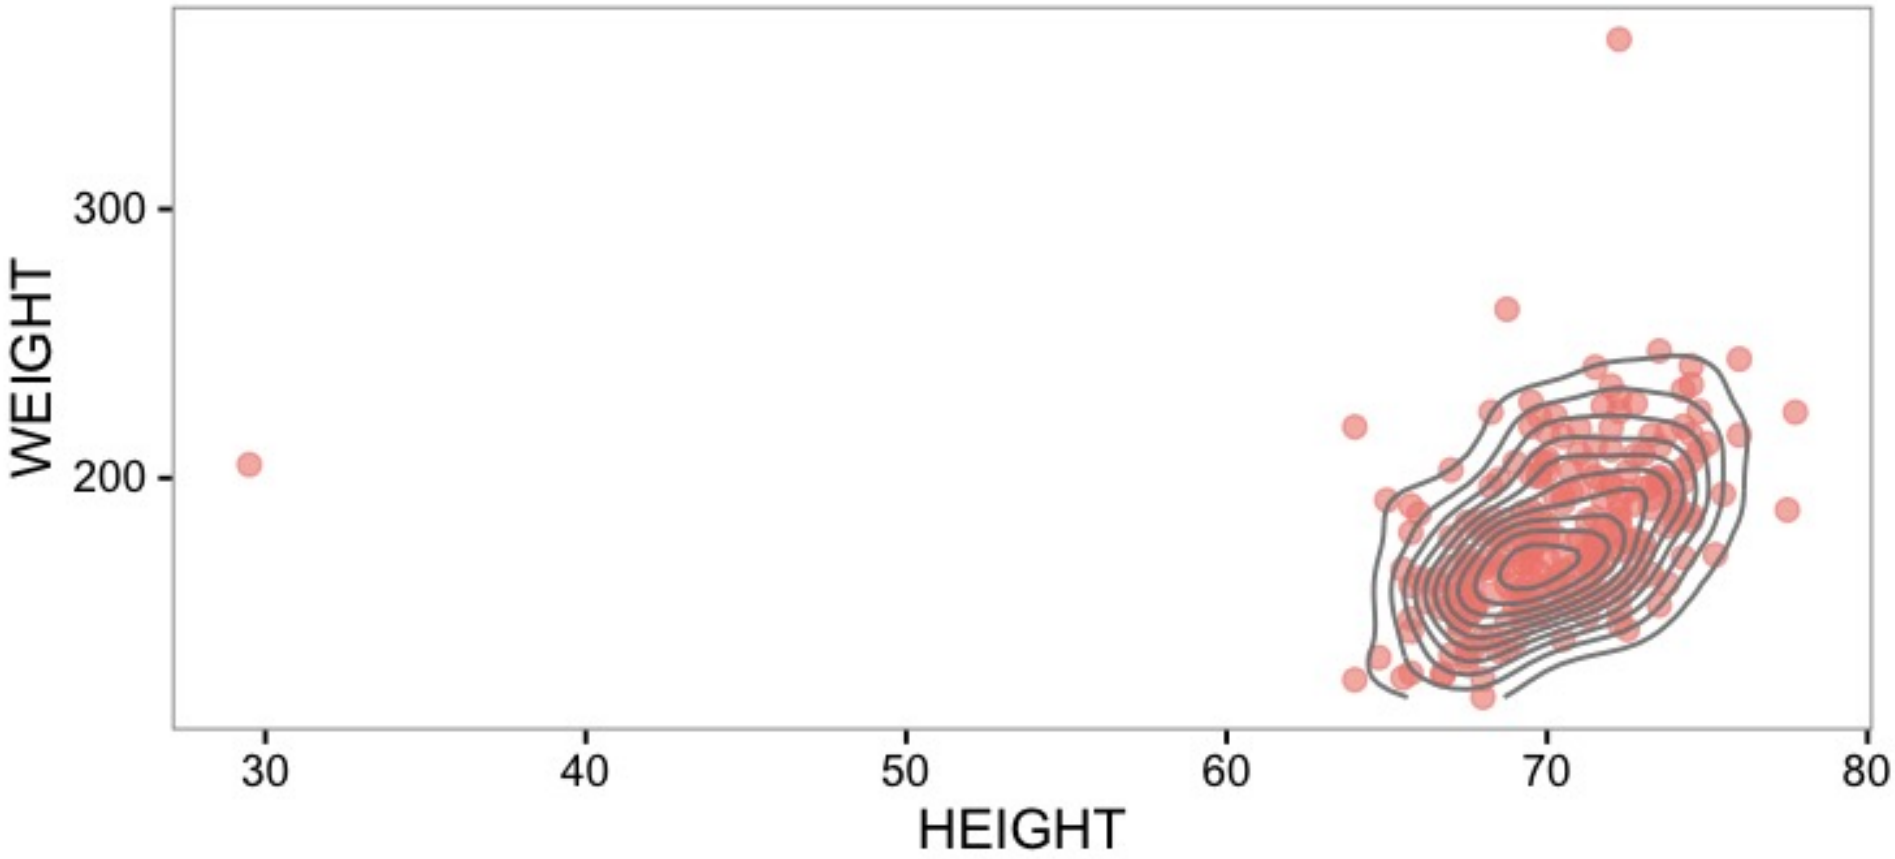

Scatter plot

- Removed of outliers & standardized

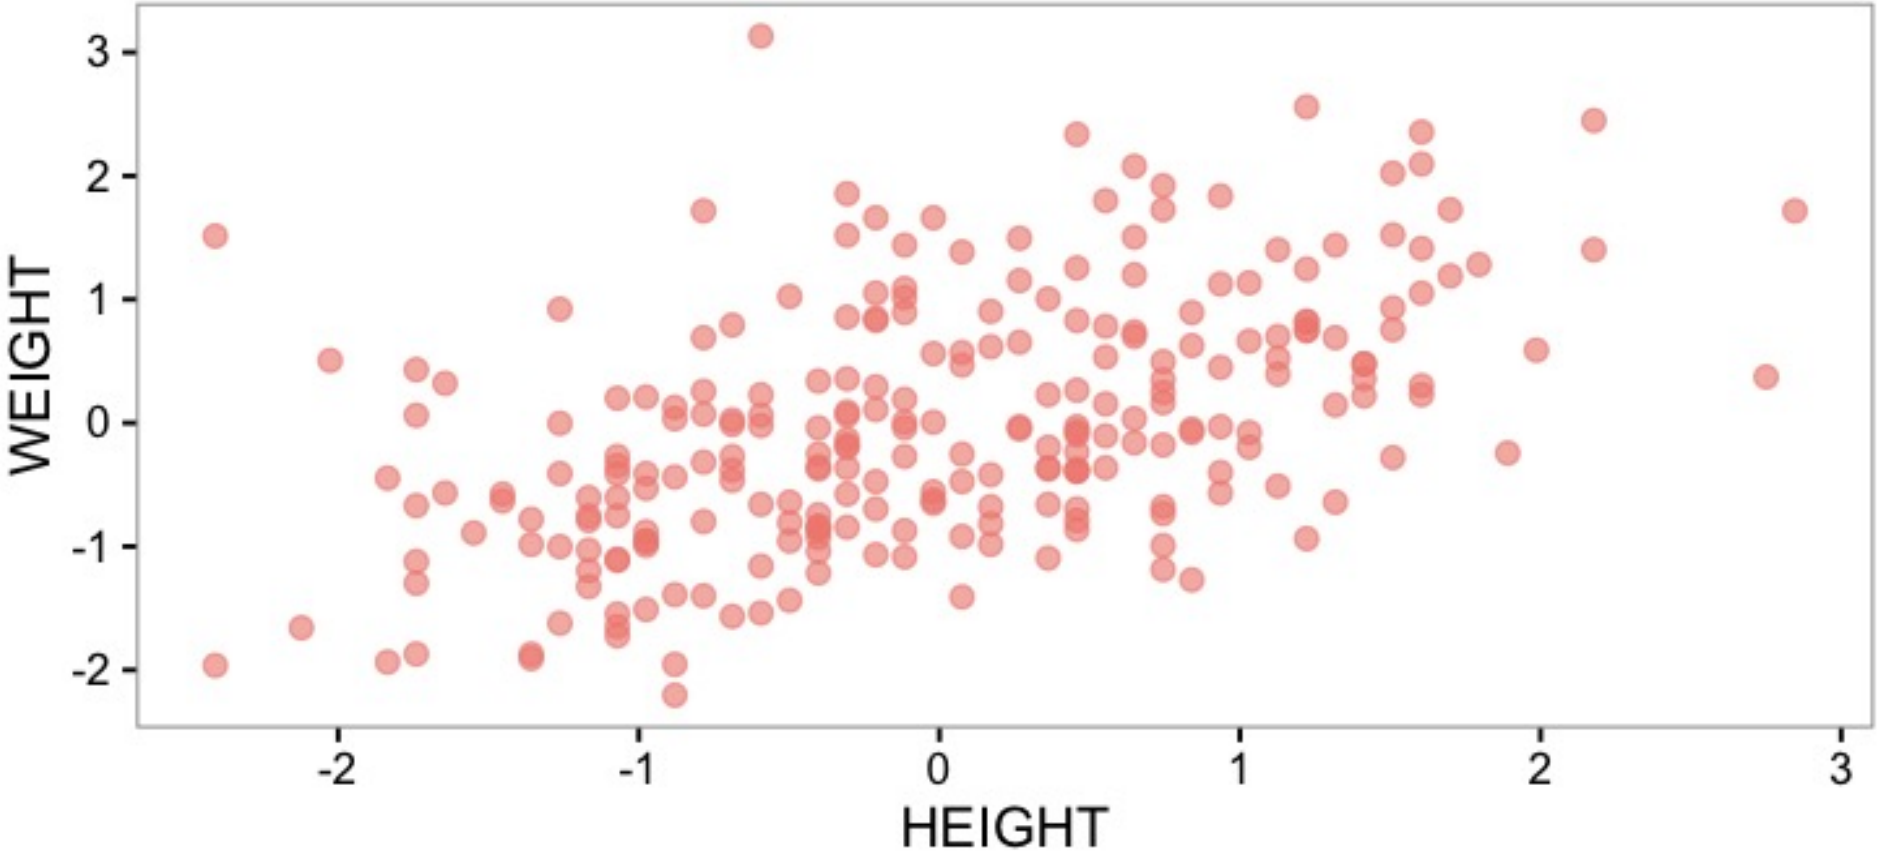

Correlation seen from scatter plots

Zero Correlation

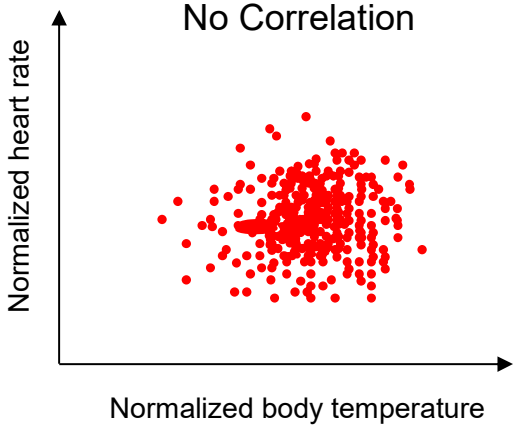

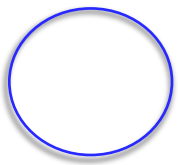

Positive correlation

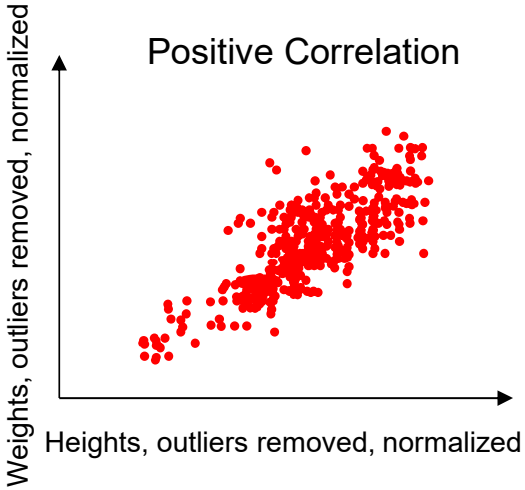

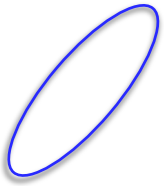

Negative correlation

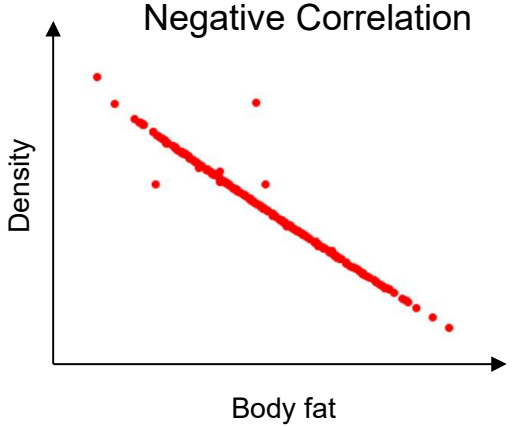

What kind of Correlation?

- Line of code in a database and number of bugs
- Frequency of hand washing and number of germs on your hands
- GPA and hours spent playing video games
- earnings and happiness

Credit: Prof. David Varodayan

Correlation Coefficient

- Given a data set $\{(x_i, y_i)\}$ consisting of items $(x_1, y_1) \ldots (x_N, y_N)$,
  - Standardize the coordinates of each feature:

$$\widehat{x_i} = \frac{x_i - mean(\{x_i\})}{std(\{x_i\})} \qquad \widehat{y_i} = \frac{y_i - mean(\{y_i\})}{std(\{y_i\})}$$

  - Define the correlation coefficient as:

$$corr(\{(x_i, y_i)\}) = \frac{1}{N} \sum_{i=1}^{N} \widehat{x_i}\widehat{y_i}$$

Correlation Coefficient

$$\widehat{x_i} = \frac{x_i - mean(\{x_i\})}{std(\{x_i\})} \qquad \widehat{y_i} = \frac{y_i - mean(\{y_i\})}{std(\{y_i\})}$$

$$corr(\{(x_i, y_i)\}) = \frac{1}{N} \sum_{i=1}^{N} \widehat{x_i}\widehat{y_i}$$

$$= mean(\{\widehat{x_i}\widehat{y_i}\})$$

Q: Correlation Coefficient

- Which of the following describe(s) correlation coefficient correctly?

  A. It's unitless

  B. It's defined in standard coordinates

  C. Both A & B

$$corr(\{(x_i, y_i)\}) = \frac{1}{N}\sum_{i=1}^{N} \hat{x}_i \hat{y}_i$$

A visualization of correlation coefficient

https://rpsychologist.com/d3/correlation/

In a data set $\{(x_i, y_i)\}$ consisting of items

(x_1, y_1) ... (x_N, y_N),

$corr(\{(x_i, y_i)\}) > 0$ shows positive correlation

$corr(\{(x_i, y_i)\}) < 0$ shows negative correlation

$corr(\{(x_i, y_i)\}) = 0$ shows no correlation

Correlation seen from scatter plots

Zero Correlation

corr ≈ 0

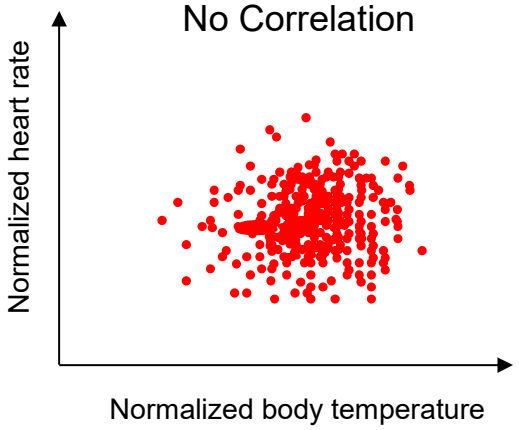

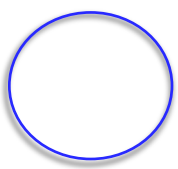

Positive correlation

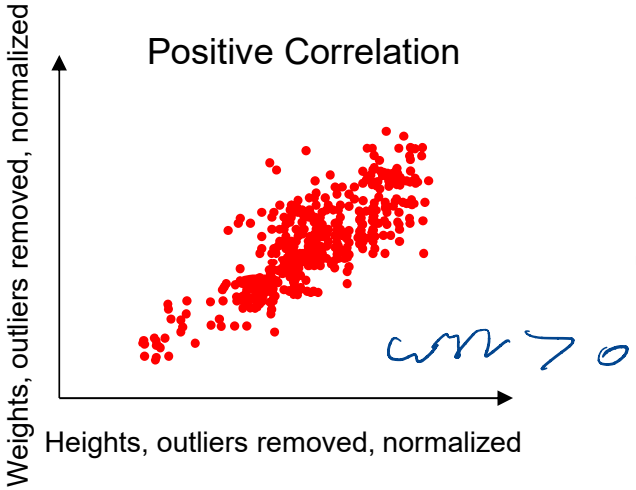

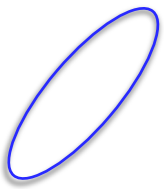

Negative correlation

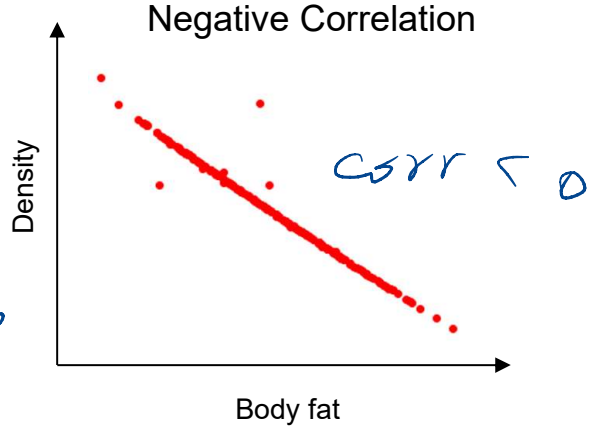

The Properties of Correlation Coefficient

- The correlation coefficient is symmetric

$$corr(\{(x_i, y_i)\}) = corr(\{(y_i, x_i)\})$$

- Translating the data does **NOT** change the correlation coefficient

The Properties of Correlation Coefficient

- Scaling the data may change the sign of the correlation coefficient

$$corr(\{(a\ x_i + b,\ \ c\ y_i + d)\})$$
$$= sign(a\ c) corr(\{(x_i, y_i)\})$$

The Properties of Correlation Coefficient

- The correlation coefficient is bounded within [-1, 1]

$corr(\{(x_i, y_i)\}) = 1$ if and only if $\widehat{x_i} = \widehat{y_i}$

$corr(\{(x_i, y_i)\}) = -1$ if and only if $\widehat{x_i} = -\widehat{y_i}$

Q. Which of the following has correlation coefficient equal to 1?

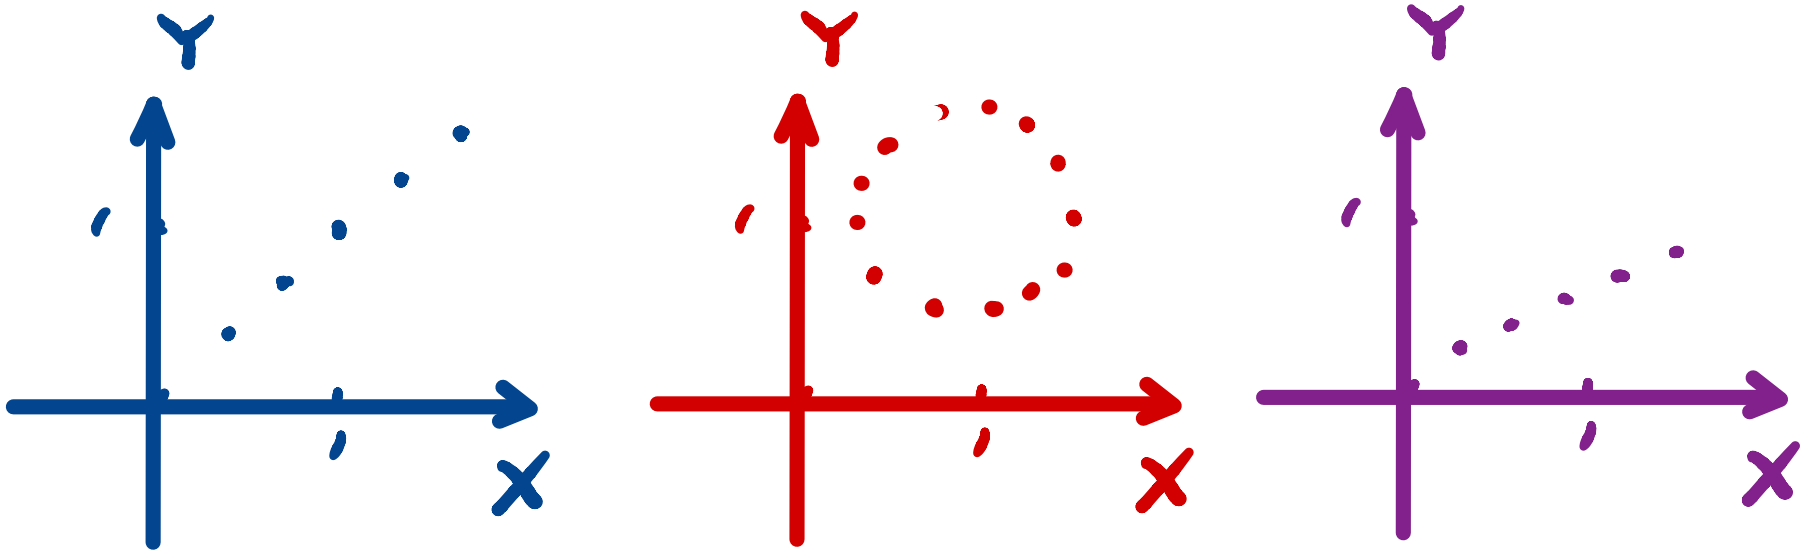

A. Left and right

B. Left

C. Middle

Review and finish it at home.

$$y = ax; \quad a > 0$$

$$\hat{y} = \frac{ax - mean(\{y\})}{std(\{y\})}$$

$$= \frac{ax - a \cdot mean(\{x\})}{a \cdot std(\{x\})}$$

$$= \hat{x}$$

$$Corr = \frac{\sum \hat{x}_i \hat{y}_i}{N}$$

Concept of Correlation Coefficient's bound

- The correlation coefficient can be written as

$$corr(\{(x_i, y_i)\}) = \frac{1}{N}\sum_{i=1}^{N} \widehat{x_i}\widehat{y_i}$$

$$corr(\{(x_i, y_i)\}) = \sum_{i=1}^{N} \frac{\widehat{x_i}}{\sqrt{N}}\frac{\widehat{y_i}}{\sqrt{N}}$$

- It's the inner product of two vectors

$$\left\langle \frac{\widehat{x_1}}{\sqrt{N}}, \quad \dots \quad \frac{\widehat{x_N}}{\sqrt{N}} \right\rangle \text{ and } \left\langle \frac{\widehat{y_1}}{\sqrt{N}}, \quad \dots \quad \frac{\widehat{y_N}}{\sqrt{N}} \right\rangle$$

Inner product

- Inner product's geometric meaning:

$$|\nu_1|\,|\nu_2|\,cos(\theta)$$

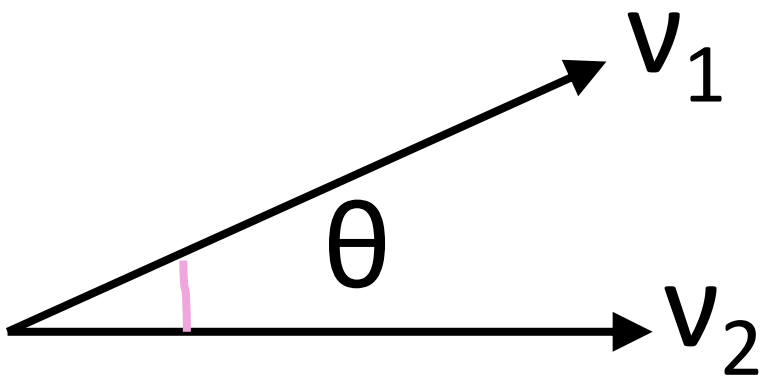

- Lengths of both vectors

$\nu_1 = \left\langle \frac{\widehat{x_1}}{\sqrt{N}}, \quad \dots \quad \frac{\widehat{x_N}}{\sqrt{N}} \right\rangle \qquad \nu_2 = \left\langle \frac{\widehat{y_1}}{\sqrt{N}}, \quad \dots \quad \frac{\widehat{y_N}}{\sqrt{N}} \right\rangle$

are 1

Bound of correlation coefficient

$$|corr(\{(x_i, y_i)\})| = |cos(\theta)| \leq 1$$

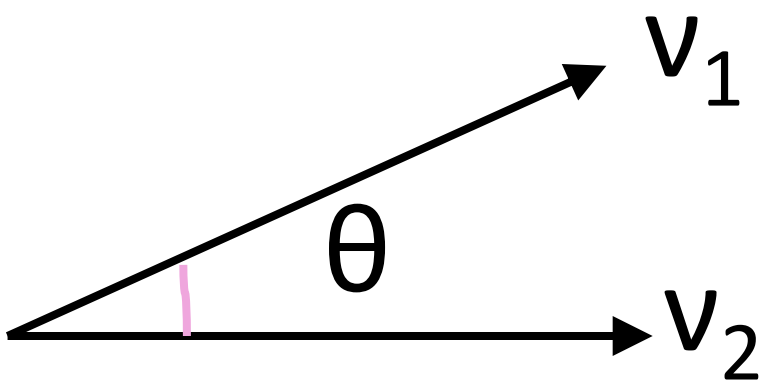

$$\nu_1 = \left\langle \frac{\widehat{x_1}}{\sqrt{N}}, \quad \cdots \quad \frac{\widehat{x_N}}{\sqrt{N}} \right\rangle \quad \nu_2 = \left\langle \frac{\widehat{y_1}}{\sqrt{N}}, \quad \cdots \quad \frac{\widehat{y_N}}{\sqrt{N}} \right\rangle$$

The Properties of Correlation Coefficient

- Symmetric
- Translating invariant
- Scaling only may change sign
- bounded within [-1, 1]

Using correlation to predict

- **Caution**! Correlation is **NOT** Causation

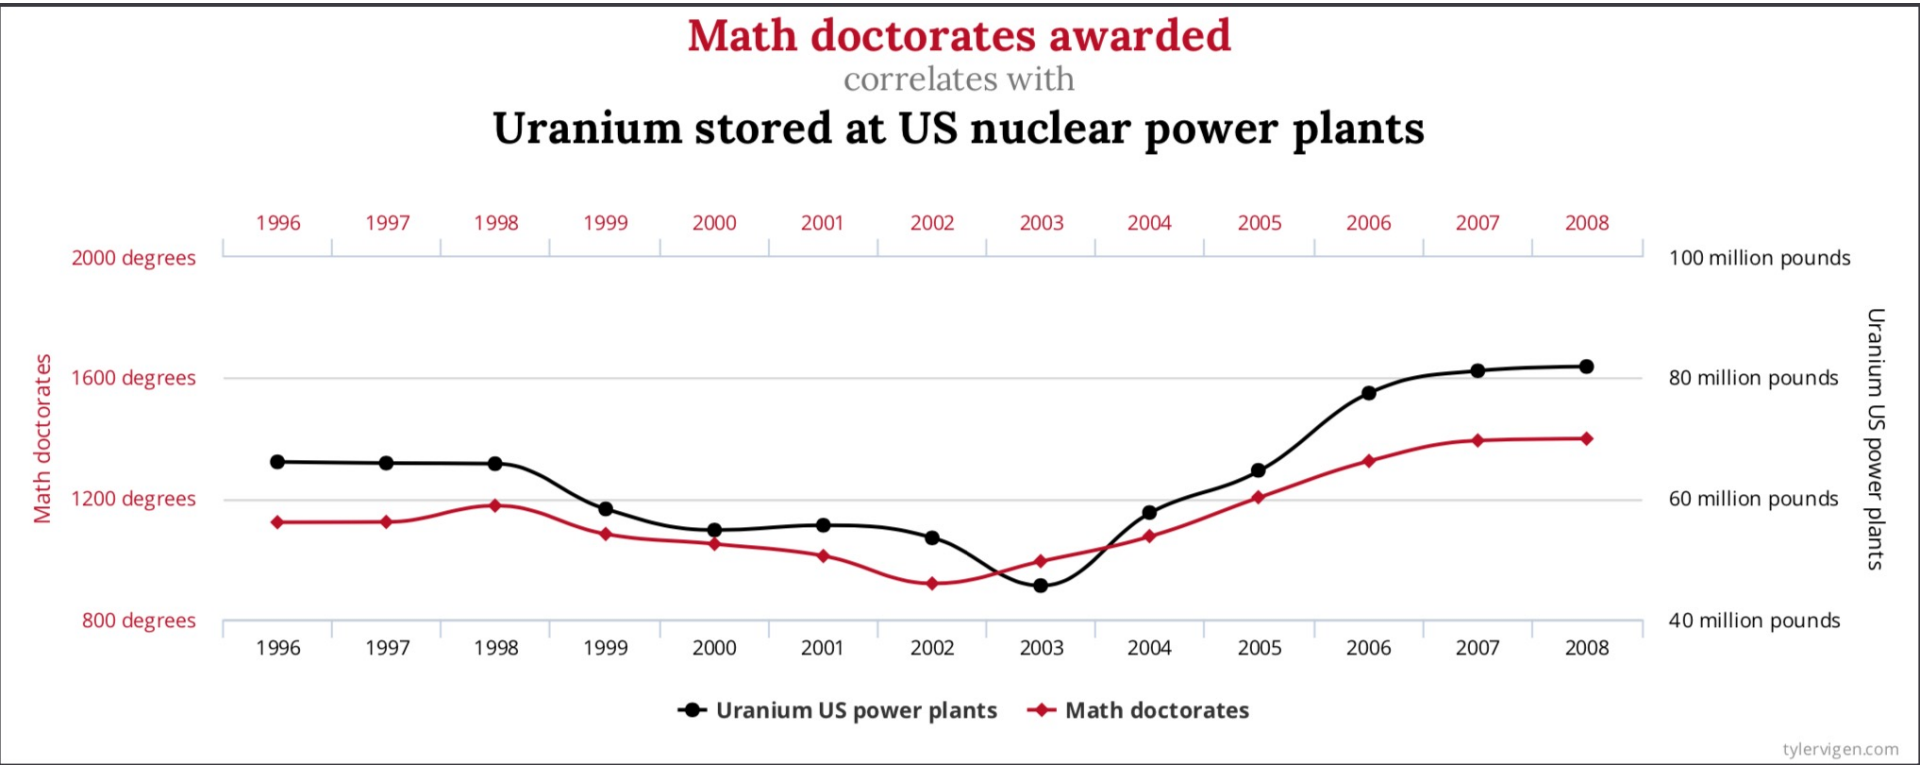

Credit: Tyler Vigen

How do we go about the prediction?

- Removed of outliers & standardized

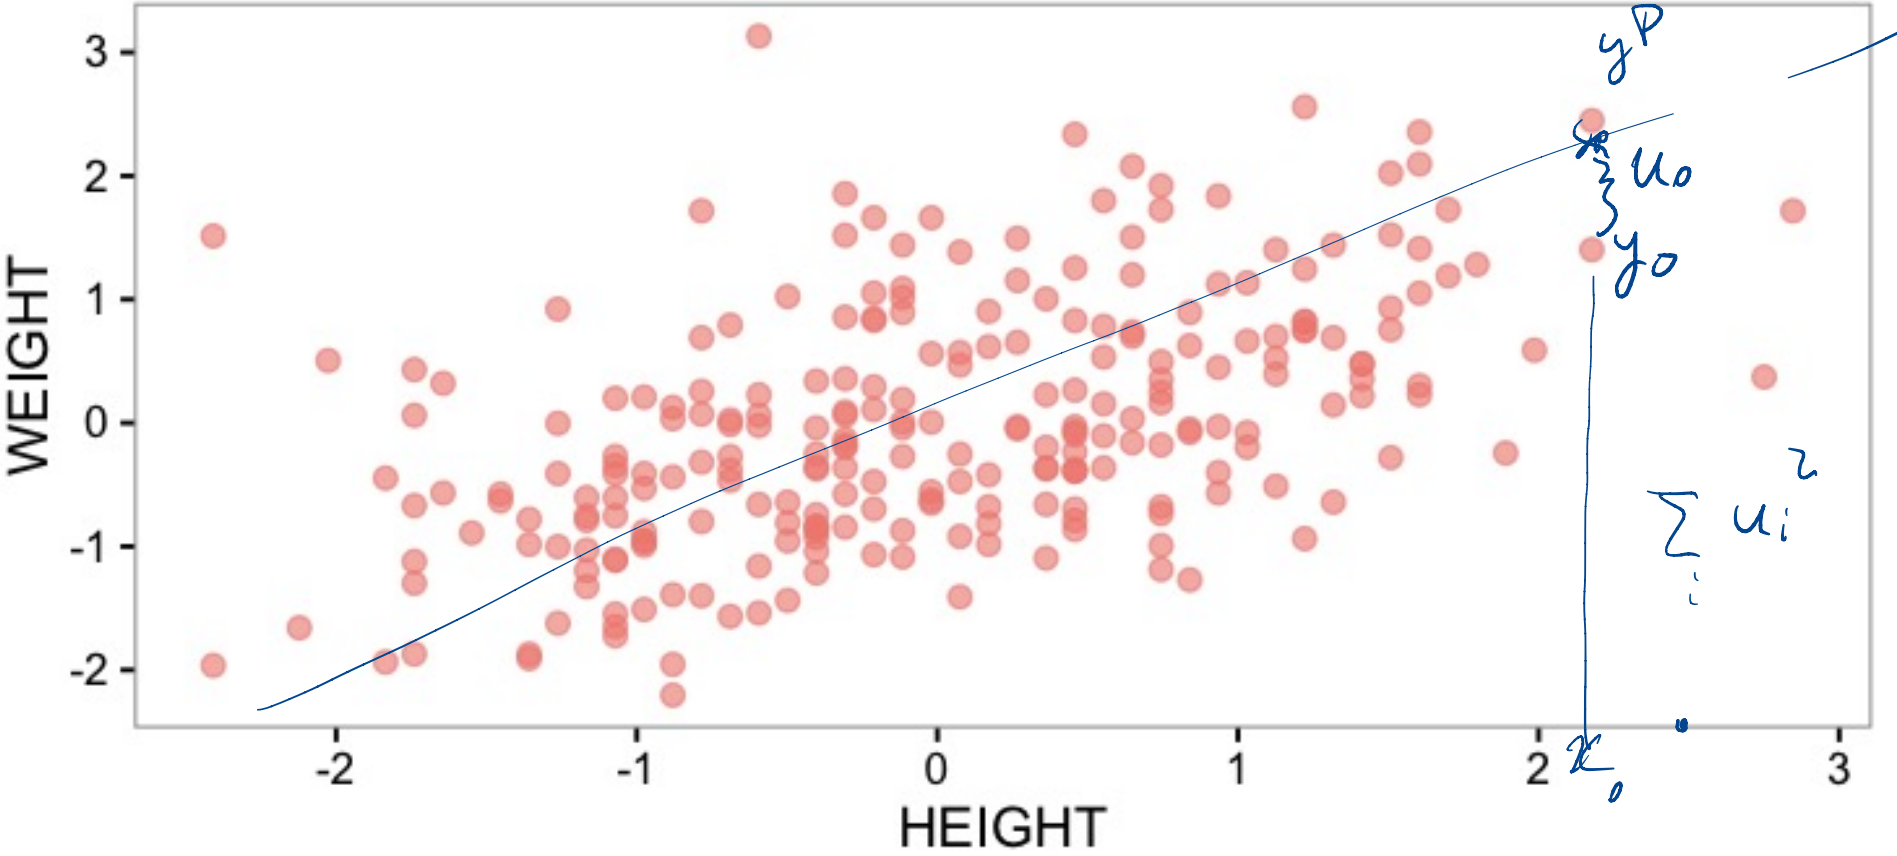

Using correlation to predict

- Given a correlated data set $\{(x_i, y_i)\}$

  we can predict a value ${y_0}^p$ that goes with x_0 a value

- In standard coordinates $\{(\widehat{x_i}, \widehat{y_i})\}$

  we can predict a value ${\widehat{y_0}}^p$ that goes with $\widehat{x_0}$ a value

Q:

- Which coordinates will you use for the predictor using correlation?

A. Standard coordinates

B. Original coordinates

C. Either

Linear predictor and its error

- We will assume that our predictor is linear

$$\widehat{y}^{\,p} = a\,\widehat{x} + b$$

- We denote the prediction at each $\widehat{x_i}$ in the data set as $\widehat{y_i}^{\,p}$

$$\widehat{y_i}^{\,p} = a\,\widehat{x_i} + b$$

- The error in the prediction is denoted u_i

$$u_i = \widehat{y_i} - \widehat{y_i}^{\,p} = \widehat{y_i} - a\,\widehat{x_i} - b$$

Require the mean of error to be zero

We would try to make the mean of error equal to zero so that it is also centered around 0 as the standardized data:

Require the variance of error is minimal

Require the mean of error to be zero

We would try to make the mean of error equal to zero so that it is also centered around 0 as the standardized data: $\text{mean}(\{u_i\}) = 0$ ✓

$$\text{mean}(\{u_i\}) = \text{mean}(\{\hat{y} - \hat{y}^p\})$$

$$= \text{mean}(\{\hat{y} - a\hat{x} - b\})$$

$$= \text{mean}(\{\hat{y}\}) - a\,\text{mean}(\{\hat{x}\})$$

($\text{mean}(\{\hat{y}\}) \to 0$, $\text{mean}(\{\hat{x}\}) \to 0$)

$$-b = 0$$

$$b = 0$$

Recall:

$$\text{mean}(\{kx+b\}) = k\,\text{mean}(\{x\}) + b$$

Require the variance of error is minimal

minimize $\mathrm{var}(\{u_i\})$

$$\mathrm{var}(\{u_i\}) = \mathrm{mean}(\{u_i - \mathrm{mean}(\{u_i\})^2\})$$

$$= \mathrm{mean}(\{u_i^2\})$$

$$= \mathrm{mean}(\{(\hat{y} - \hat{y}^p)^2\}) \qquad u_i = \hat{y} - \hat{y}^p$$

$$= \mathrm{mean}(\{(\hat{y} - a\hat{x})^2\}) \qquad = \hat{y} - a\hat{x}$$

$$= \mathrm{mean}(\{\hat{y}^2 - 2a\hat{x}\hat{y} + a^2\hat{x}^2\}) \qquad \because b = 0$$

$$= \mathrm{mean}(\{\hat{y}^2\}) - 2a\,\mathrm{mean}(\{\hat{x}\hat{y}\}) + a^2 \cdot \mathrm{mean}(\{\hat{x}^2\})$$

Require the variance of error is minimal

minimize $var(\{u_i\})$

$$var(\{u_i\}) = mean(\{u_i - mean(\{u_i\})^2\})$$

$$= mean(\{u_i^2\})$$

$$= mean(\{(\hat{y} - \hat{y}^p)^2\})$$

$u_i = \hat{y} - \hat{y}^p = \hat{y} - a\hat{x}$ $\because b = 0$

$$= mean(\{(\hat{y} - a\hat{x})^2\})$$

$$= mean(\{\hat{y}^2 - 2a\hat{x}\hat{y} + a^2\hat{x}^2\})$$

$$= mean(\{\hat{y}^2\}) - 2a\,mean(\{\hat{x}\hat{y}\}) + a^2 mean(\{\hat{x}^2\})$$

$mean(\{\hat{y}^2\})$ $= 1$

$= mean(\{(\hat{y} - 0)^2\})$

$= mean(\{(\hat{y} - mean(\hat{y}))^2\})$

$= var(\{\hat{y}\}) = 1$

$mean(\{\hat{x}\hat{y}\}) \rightarrow corr \rightarrow r$

Require the variance of error is minimal

$$\mathrm{var}(\{u\}) = \mathrm{mean}(\{\hat{y}^2\}) - 2a\,\mathrm{mean}(\{\hat{x}\,\hat{y}\}) + a^2\,\mathrm{mean}(\{\hat{x}^2\})$$

$$= 1 - 2a\,\mathrm{mean}(\{\hat{x}\,\hat{y}\}) + a^2$$

$$= 1 - 2a\,\mathrm{corr}(\{x, y\}) + a^2$$

$$r = \mathrm{corr}(\{x, y\})$$

$$= 1 - 2ar + a^2$$

$\mathrm{var}(\{u\})$ ↓ argmin

$a = ?$

$$\frac{d\,\mathrm{var}(\{u\})}{da} = -2r + 2a = 0$$

Require the variance of error is minimal

$$\hat{y}^p = a\hat{x} + b$$

$$= r\hat{x}$$

$$a = r$$

$$b = 0$$

Here is the linear predictor!

$$\widehat{y}^{\,p} = r\ \widehat{x}$$

Correlation coefficient

Prediction Formula

- In standard coordinates

$$\widehat{y_0}^{\,p} = r\ \widehat{x_0} \quad \text{where} \quad r = corr(\{(x_i, y_i)\})$$

- In original coordinates

$$\frac{y_0^p - mean(\{y_i\})}{std(\{y_i\})} = r\frac{x_0 - mean(\{x_i\})}{std(\{x_i\})}$$

Root-mean-square (RMS) prediction error

- $Given \quad var(\{u_i\}) = 1 - 2ar + a^2$
  $\&$ $a = r$

$var(\{u_i\}) = 1 - r^2$

- $$RMS \;\; error = \sqrt{mean(\{u_i^2\})}$$
  $$= \sqrt{var(\{u_i\})}$$
  $$= \sqrt{1 - r^2}$$

See the error through simulation

https://rpsychologist.com/d3/correlation/

Example: Body Fat data

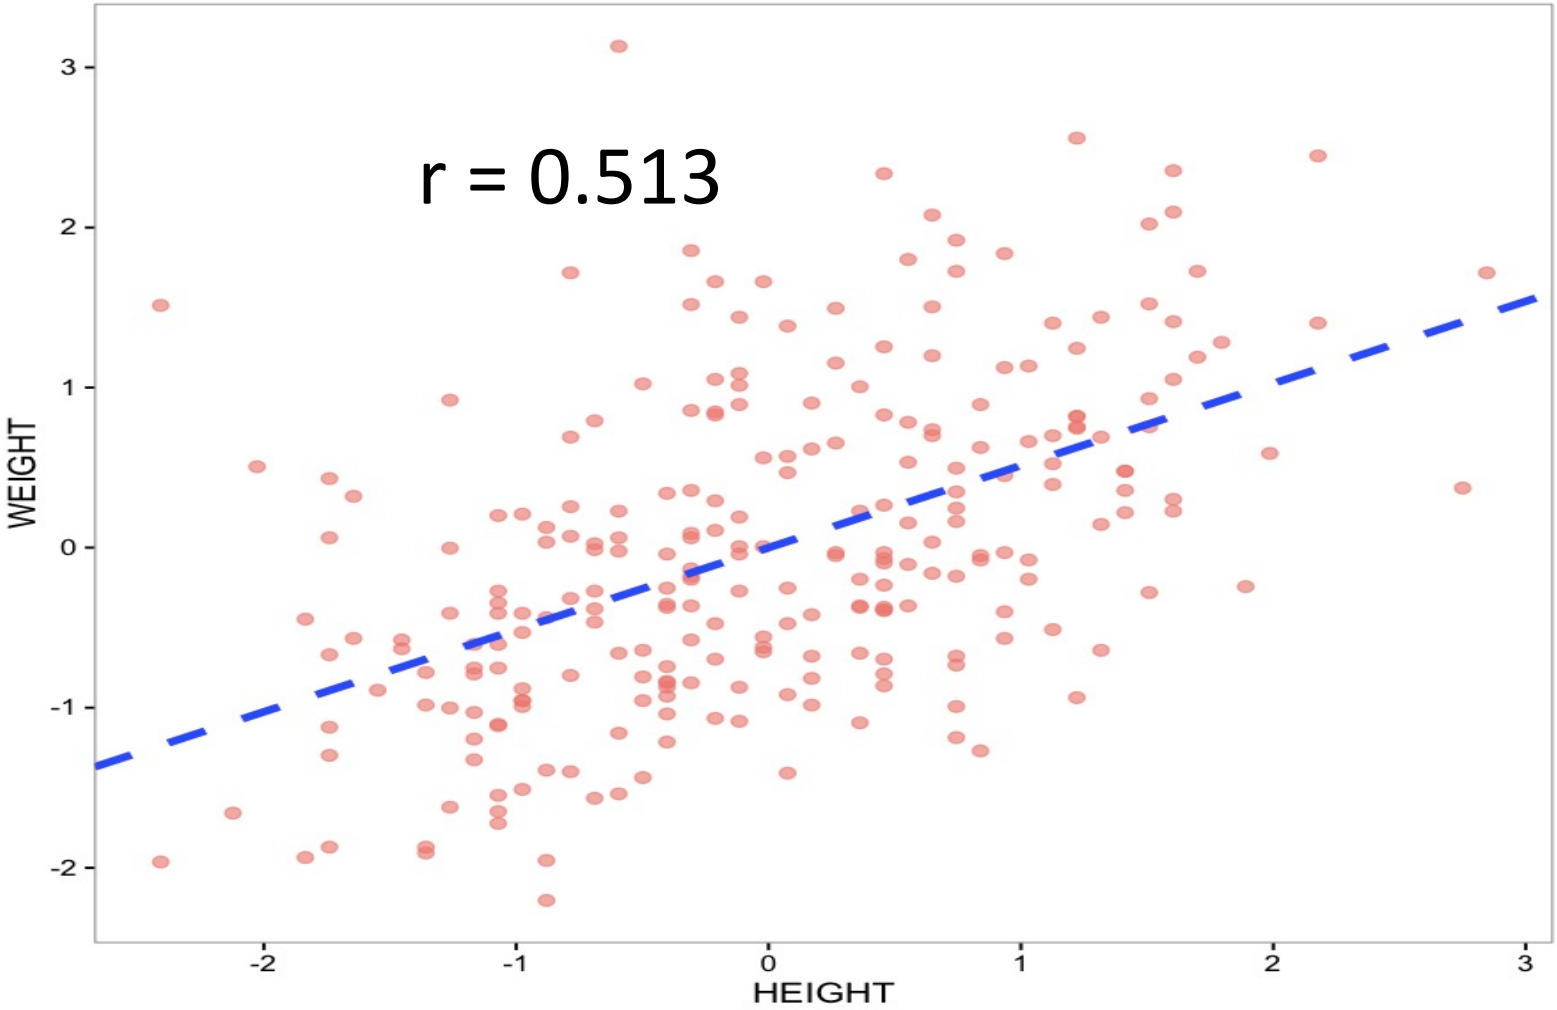

Example: remove 2 more outliers

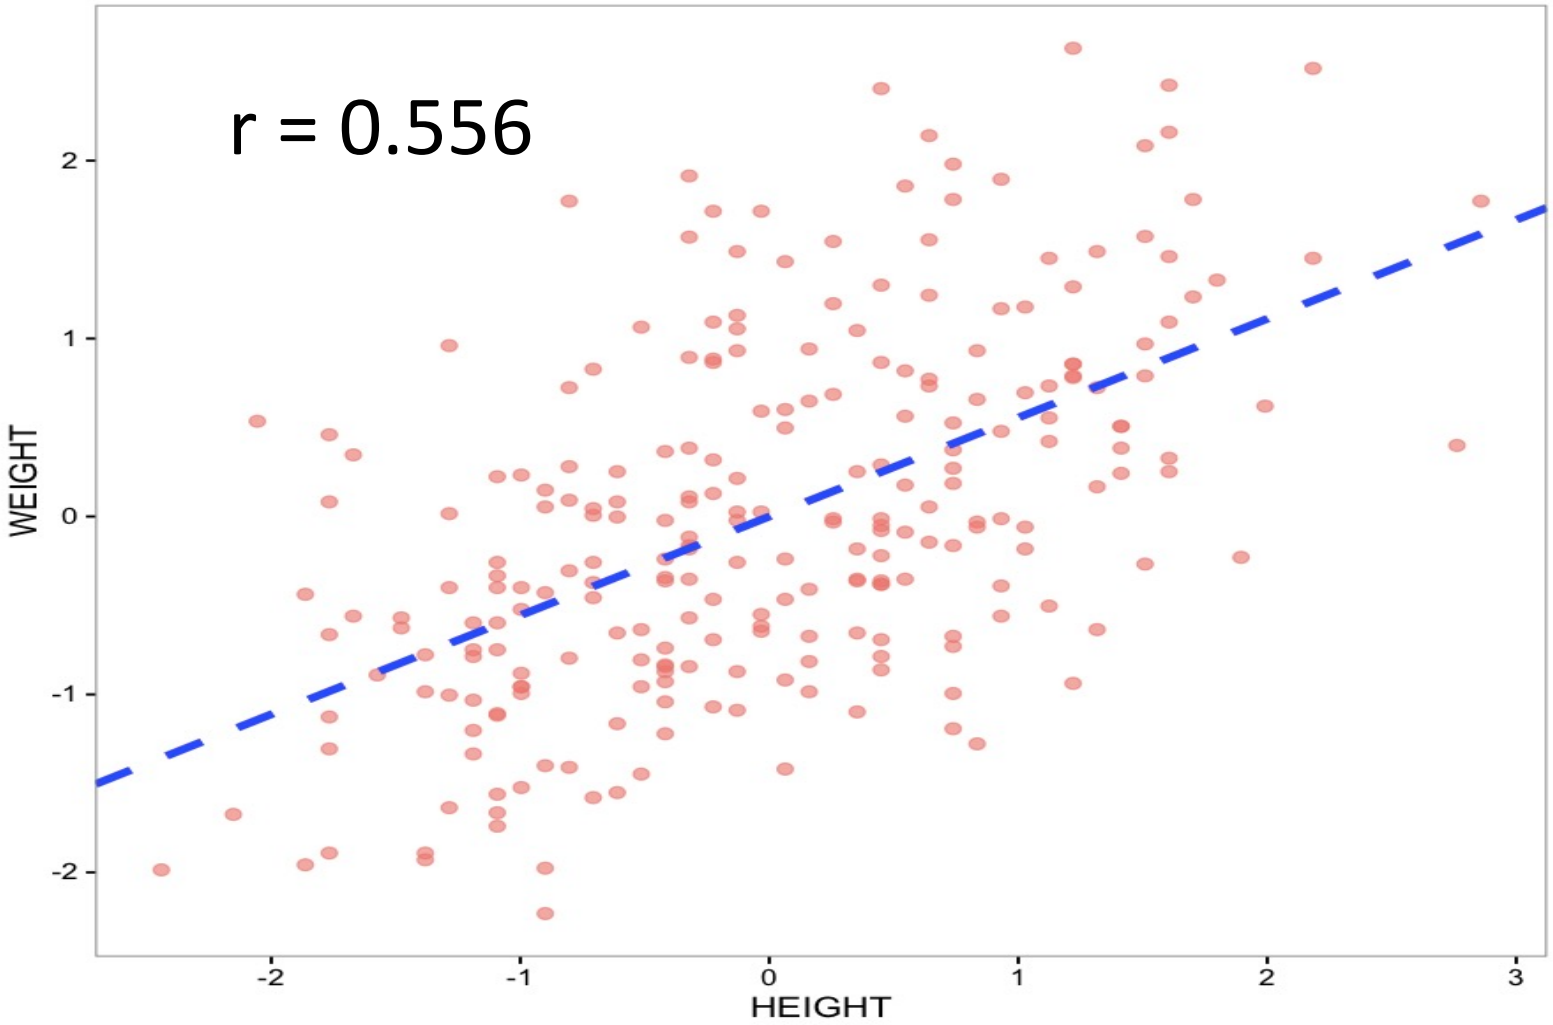

Scatter plot

- Coupled with heatmap to show a 3rd feature

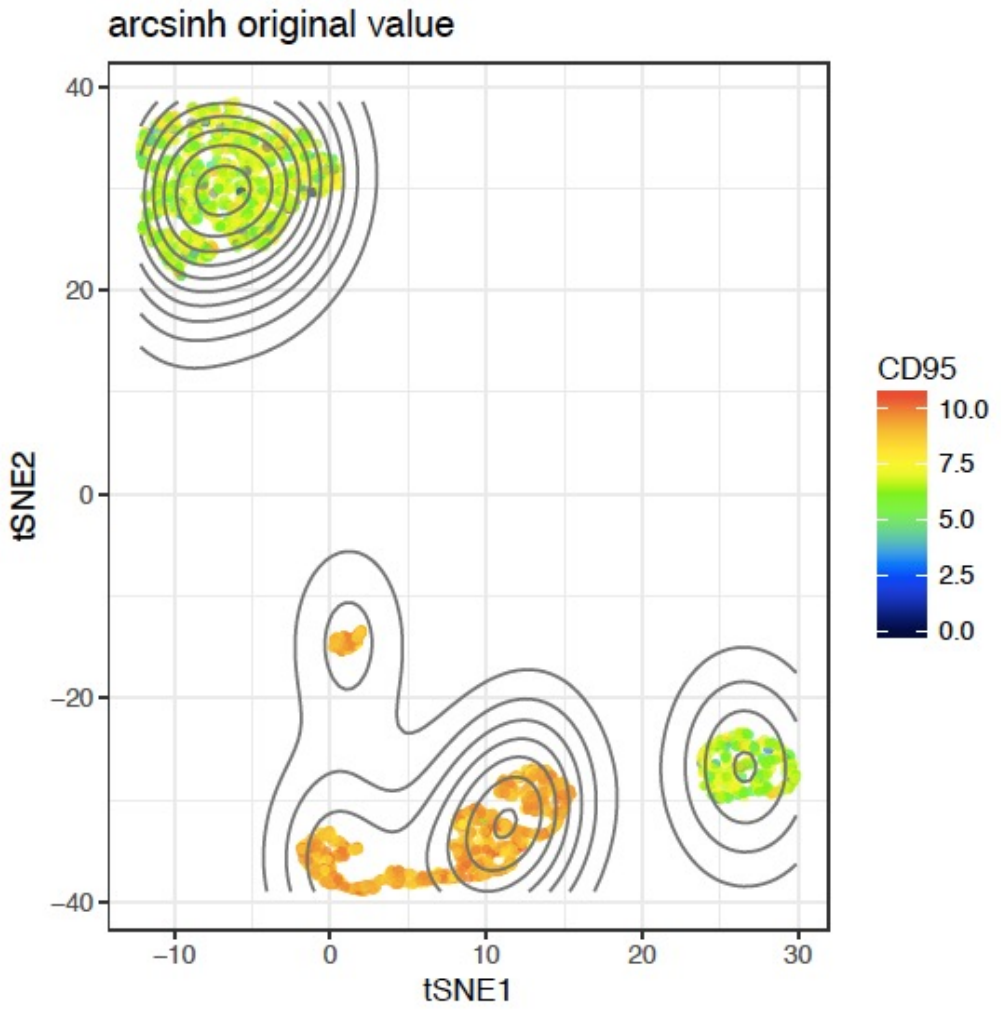

Time Series Plot: Stock of Amazon

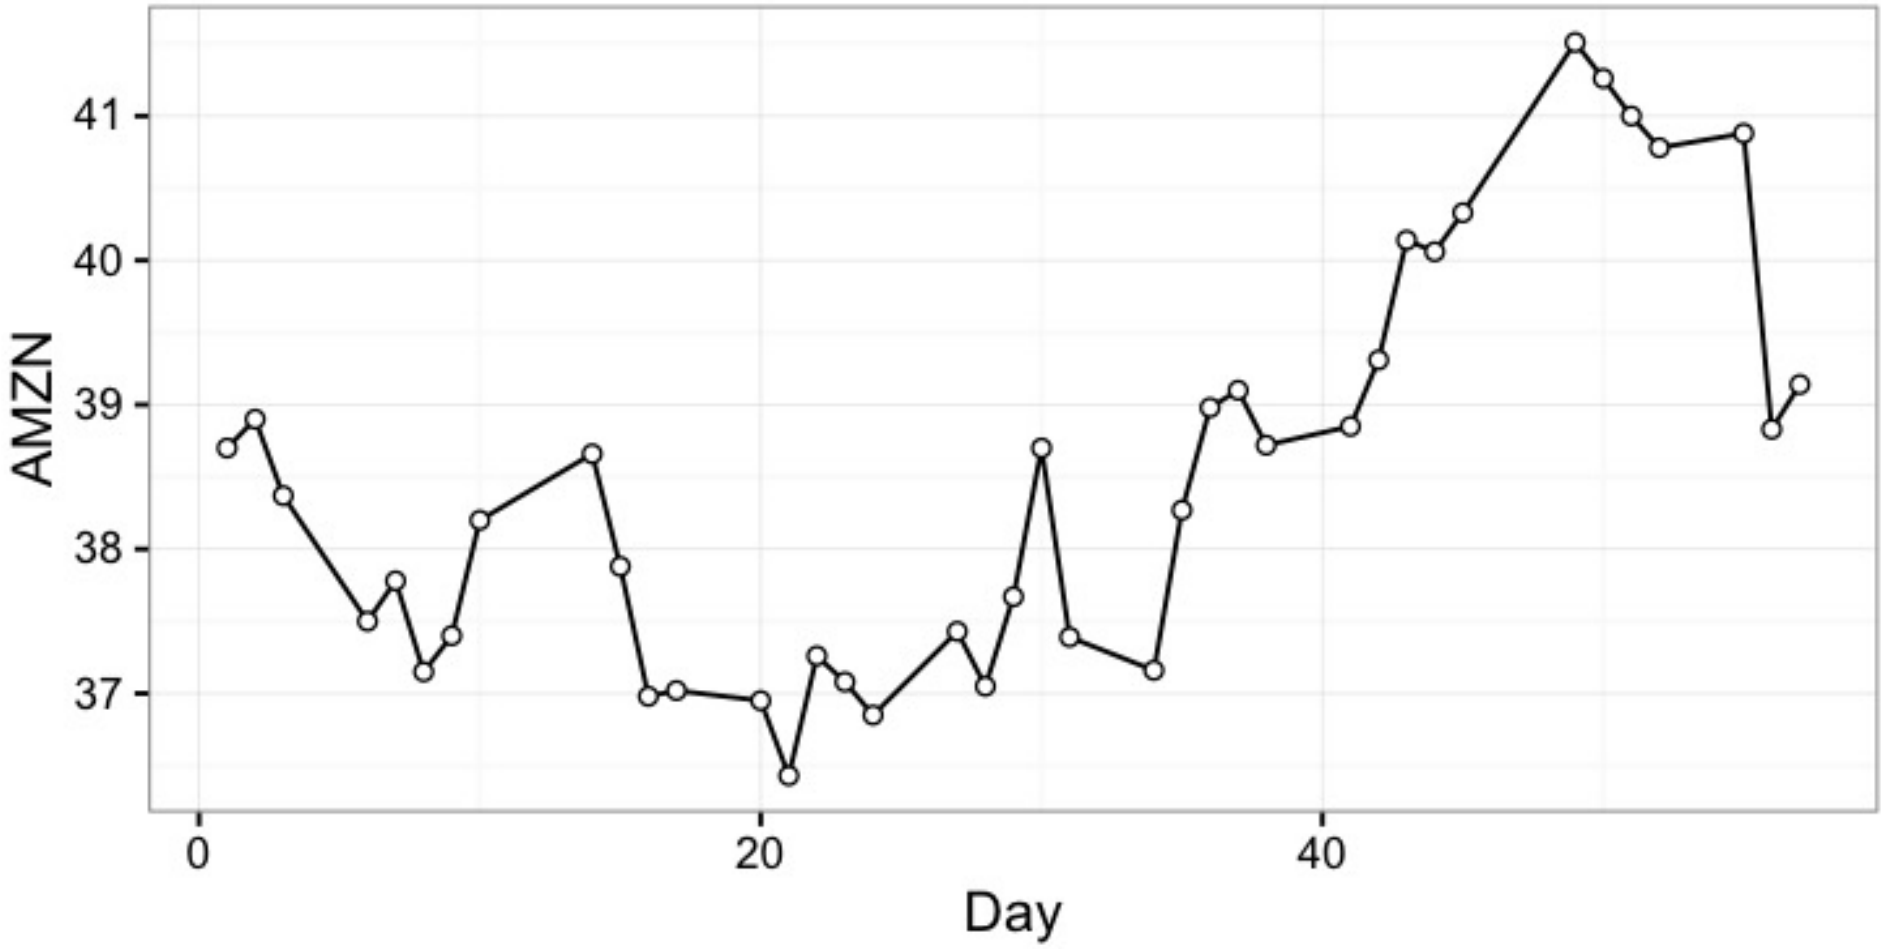

Heatmap

- Display matrix of data via gradient of color(s)

Summarization of 4 locations' annual mean temperature by month

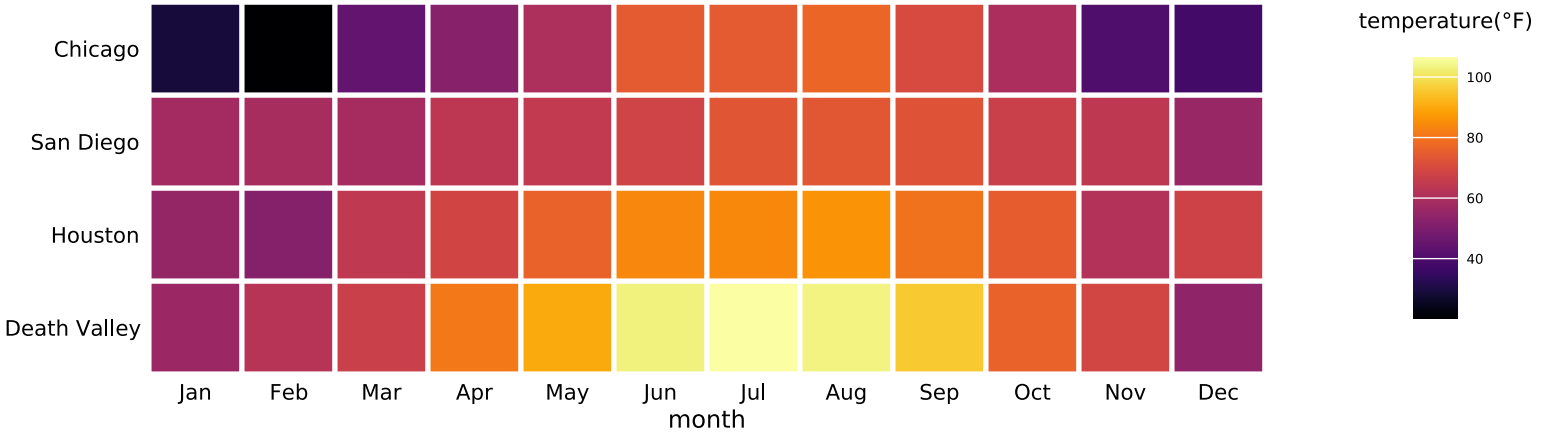

3D bar chart

- Transparent 3D bar chart is good for small # of samples across categories

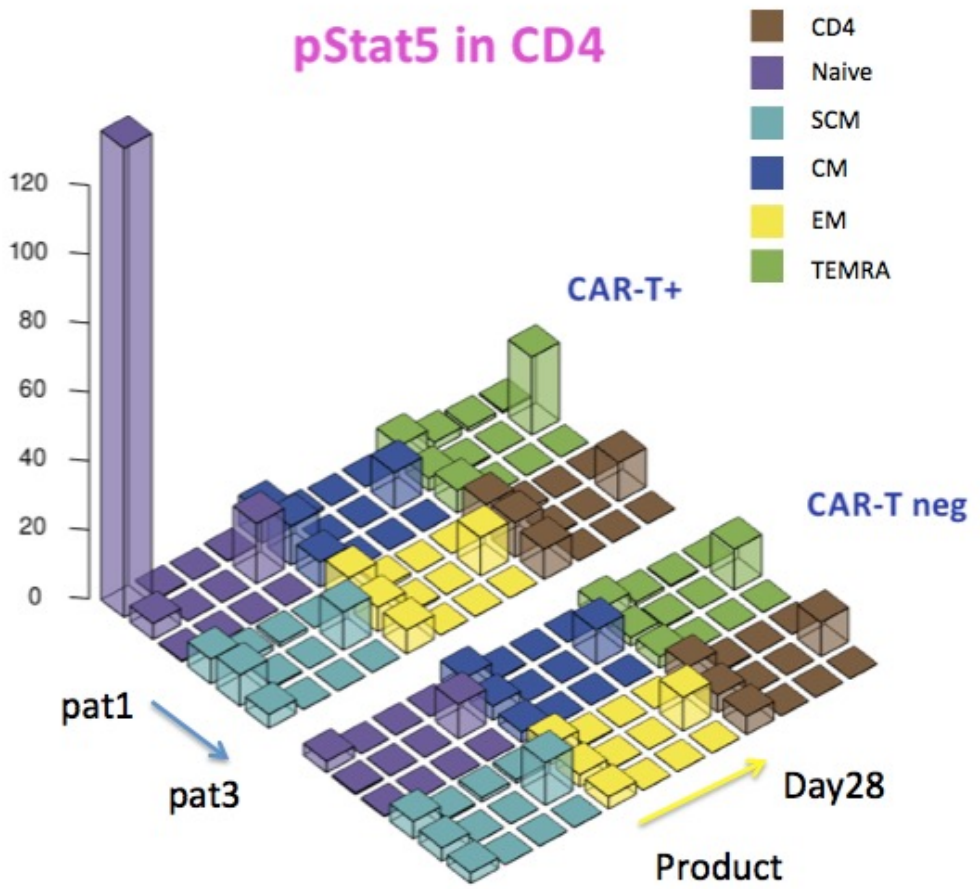

Assignments

- Quiz1 open at 4:30pm today on PL
- Finish reading Chapter 2 of the textbook
- Work on the Week 2 module on Canvas
- Next time: Probability a first look

Additional References

- Charles M. Grinstead and J. Laurie Snell "Introduction to Probability”
- Morris H. Degroot and Mark J. Schervish "Probability and Statistics”

See you next time

See You!